

The Three Pillars of LLM Success

Data, Algorithms, and Hardware

Why Did LLMs Succeed Now?

The ideas behind neural networks have existed since the 1950s.

- So why did LLMs like GPT and Claude only become possible in recent years?

Because **three factors** finally converged at the same time: DHA — Data, Hardware, and Algorithms.

Data	Hardware	Algorithms
Massive text from the internet	GPUs powerful enough to train on it all	The Transformer architecture (2017)

Remove any one of these, and modern LLMs would not exist.

The Three Pillars — An Analogy

Think of building a skyscraper:

Pillar	Skyscraper Analogy	LLM Equivalent
Data	Raw materials (steel, concrete, glass)	Trillions of tokens from books, web, code
Hardware	Construction equipment (cranes, trucks)	GPUs, TPUs, large compute clusters
Algorithms	Architectural blueprint	Transformer, attention mechanism

You can have the best blueprint in the world, but without materials and equipment, nothing gets built. The same is true for LLMs.

Pillar 1: Data

Data — The Fuel for LLMs

LLMs learn language by reading **enormous amounts of text**. The quality, quantity, and diversity of data determine how capable the model becomes.

GPT-3 (2020): trained on ~300 billion tokens

LLaMA 2 (2023): trained on ~2 trillion tokens

LLaMA 3 (2024): trained on ~15 trillion tokens

For reference: 1 trillion tokens ≈ roughly 750 billion words — that is approximately **7.5 million copies** of the entire Harry Potter series.

Where Does the Training Data Come From?

Source	Examples	What the LLM Learns
Web pages	Common Crawl, Wikipedia	General knowledge, facts
Books	Project Gutenberg, digitized libraries	Long-form reasoning, narrative
Code	GitHub, Stack Overflow	Programming languages, logic
Academic papers	arXiv, PubMed	Scientific reasoning
Conversations	Reddit, forums	Dialogue, Q&A patterns

Data Quality Matters More Than Quantity

Simply having *more* data is not enough. The training pipeline involves extensive **data cleaning**:

1. **Deduplication** — Remove duplicate web pages and repeated content
2. **Filtering** — Remove low-quality, toxic, or spam content

3. Language identification — Separate English from other languages

4. Domain balancing — Ensure a healthy mix (code, science, conversation...)

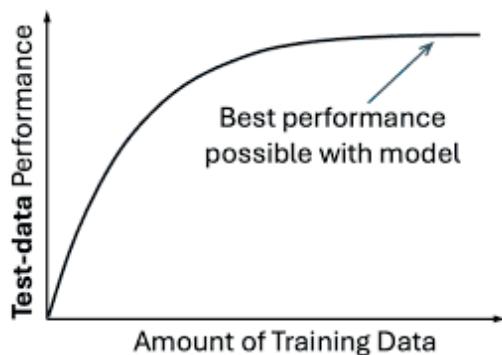
5. PII removal — Strip personal information (emails, phone numbers)

Lesson: Garbage in → garbage out. A model trained on low-quality data produces low-quality outputs, no matter how big it is.

The Data Scaling Law

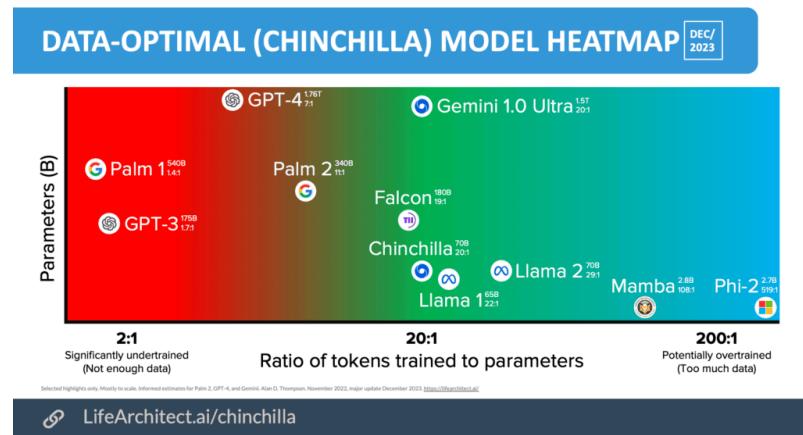
Research from OpenAI (Kaplan et al., 2020) showed a **power-law relationship**:

As you increase the amount of training data, the model's performance improves **predictably** — following a smooth curve.



Key insight from the Chinchilla paper (DeepMind, 2022):

Tokens (data) and parameters (model size) must be scaled together for optimal performance.



Tokens

- Smallest text units the model reads
- Words, subwords, or symbols

Example:

"Students learning AI"

→ ["Student", "s", "learning", "AI"]

Role: Training data

Parameters

- Learnable numbers inside the model
- Weights in Attention, FFN, Embeddings

Example:

```
W =  
[ 0.2 -1.1 0.7 0.3  
-0.8 0.1 0.6 -0.2 ]
```

Role: Knowledge storage

Learning = Adjusting parameters to minimize loss on training data

Intuition

- Tokens = Books
- Parameters = Brain capacity

More books fill the brain.

Big brain + few books → undertrained.

Question

How much data (token) is enough per parameter? how many tokens do we need?

5? 10? 100? Is more always better?

- Most early LLMs were **undertrained** — they had too many parameters for the amount of data
- The optimal ratio: **~20 tokens per parameter**
- A 70B parameter model needs ~1.4 trillion tokens to be optimally trained

This shifted the industry from "bigger models" to "**more and better data.**"

The Data Wall Problem

The AI industry is approaching a **data wall**:

- High-quality text on the internet is **finite**
- Some estimates suggest we will exhaust publicly available quality text by 2026–2028

Current strategies to overcome this:

- **Synthetic data** — use LLMs to generate training data for other LLMs
- **Data augmentation** — rephrase, translate, or restructure existing data

- **Multimodal data** — incorporate images, video, and audio
- **Proprietary data** — partnerships with publishers, news organizations
- **Curriculum learning** — train on easy data first, then harder data

This is one of the most active areas of AI research today.

Pillar 2: Hardware

Why Hardware Matters

Training an LLM is an **extraordinarily compute-intensive** task.

GPT-3 training (2020):

- ~3,640 petaflop-days of compute
- Estimated cost: **\$4.6 million** in GPU time

GPT-4 training (2023):

- Estimated cost: **\$100+ million**

LLaMA 3 405B (2024):

- Trained on **16,384 NVIDIA H100 GPUs** for ~54 days
- Estimated cost: **\$200+ million**

Without modern GPU/TPU hardware, training these models would take **centuries** on traditional CPUs.

CPU vs GPU – Why GPUs?

CPU

- Few powerful cores (8–64)
- Great for **sequential** tasks
- General purpose
- Analogy: A **few expert chefs** cooking one dish at a time

GPU

- Thousands of smaller cores (10,000+)
- Great for **parallel** tasks
- Optimized for matrix math
- Analogy: **Thousands of cooks** each doing one simple step simultaneously

LLM training is **matrix multiplication at massive scale** — exactly what GPUs excel at.

The NVIDIA Dominance

NVIDIA GPUs have become the backbone of AI training:

GPU	Year	Memory	AI Performance
V100	2017	32 GB	125 TFLOPS (FP16)
A100	2020	80 GB	312 TFLOPS (FP16)
H100	2022	80 GB	990 TFLOPS (FP16)
H200	2024	141 GB	990 TFLOPS (FP16)
B200	2024	192 GB	2,250 TFLOPS (FP16)

Key insight: Each generation roughly **3x the performance** of the previous one. This exponential improvement in hardware has directly enabled exponential growth in model size.

Beyond GPUs – Other AI Hardware

Hardware	Company	Key Feature
TPU (Tensor Processing Unit)	Google	Custom-designed for TensorFlow/JAX workloads
Trainium / Inferentia	AWS	Cost-optimized for training and inference
Gaudi	Intel/Habana	Alternative to NVIDIA for training
Groq LPU	Groq	Extremely fast inference (low latency)
Apple Neural Engine	Apple	On-device AI in iPhones/Macs
Custom ASICs	Meta, Microsoft	Purpose-built chips for internal workloads

The Memory Problem

A major hardware bottleneck is **GPU memory (VRAM)**.

A model with 70 billion parameters:

- At FP32 (32-bit): needs **280 GB** just to store weights
- At FP16 (16-bit): needs **140 GB**
- At INT8 (8-bit quantized): needs **70 GB**
- At INT4 (4-bit quantized): needs **35 GB**

A single H100 GPU has only 80 GB of memory. Training requires additional memory for gradients and optimizer states (often 3–4× the model size).

Solution: Distribute across **many GPUs** using parallelism strategies (data parallel, tensor parallel, pipeline parallel).

Inference vs Training Hardware

Training and inference have **different** hardware requirements:

Aspect	Training	Inference
Goal	Learn model weights	Generate responses
Compute	Extremely high	Moderate
Duration	Weeks to months	Milliseconds to seconds
GPUs needed	Thousands	One to a few
Cost focus	Total training cost	Cost per token/query
Key metric	Throughput (tokens/sec)	Latency (time to first token)

This is why you can run LLMs on your laptop
(inference) but cannot train one at home (training).

Quantization — Making LLMs Run on Consumer Hardware

Quantization reduces the precision of model weights to use less memory:

Original (FP16): 0.23456789... → stored in 16 bits

Quantized (INT4): 0.23456789... → mapped to nearest of 16 values
stored in just 4 bits

Precision	Bits per Weight	70B Model Size	Quality
FP16	16	~140 GB	Full quality
INT8	8	~70 GB	Near-full quality
INT4 (Q4)	4	~35 GB	Good quality, some loss
INT2	2	~17.5 GB	Noticeable degradation

Tools like llama.cpp and Ollama use quantization to let you run LLaMA 70B on a desktop with a 48 GB GPU or even a Mac with 64 GB RAM.

Pillar 3: Algorithms

The Key Breakthrough: The Transformer

Before 2017, the dominant models for language were **RNNs** (Recurrent Neural Networks) and **LSTMs**. They processed text **one word at a time**, left to right.

The problem: They were slow and struggled with long-range dependencies.

In 2017, Google published "**Attention Is All You Need**" – introducing the **Transformer** architecture.

The key idea: Process all words in a sentence **simultaneously** using a mechanism called **self-attention**.

RNN vs Transformer – Visual Comparison

RNN (Sequential)

"The cat sat on the mat"

The → cat → sat → on → the → mat
↓ ↓ ↓ ↓ ↓ ↓
h₁ → h₂ → h₃ → h₄ → h₅ → h₆

Each word must wait for the previous one.

Slow, hard to parallelize.

Transformer (Parallel)

"The cat sat on the mat"

The cat sat on the mat

All words attend to all other words simultaneously

All words processed at once.
Fast, highly parallelizable.

Self-Attention – The Core Idea

Self-attention lets each word "look at" every other word in the sentence to understand context.

Example: "*The **bank** of the river was covered in mud.*"

For the word "bank," self-attention computes:

- How relevant is "river"? → **Very high** (helps disambiguate meaning)
- How relevant is "mud"? → **High** (confirms "river bank" meaning)
- How relevant is "The"? → **Low**

This allows the model to understand that "bank" means **riverbank**, not a financial institution — by attending to context words.

Self-Attention — Simplified Math

Imagine a classroom

Every word is a kid in the classroom.

- They want to talk to each other to understand the sentence.
- But before talking, each kid makes 3 cards.

QKV – The Three Vectors

For each word, the model computes three vectors:

- **Q (Query)**: "What am I looking for?"
- **K (Key)**: "What do I contain?"
- **V (Value)**: "What information do I provide?"

$$\text{Attention}(Q, K, V) = \text{softmax} \left(\frac{QK^T}{\sqrt{d_k}} \right) V$$

In plain English:

1. Compare each word's Query with every other word's Key → get similarity scores
2. Normalize with softmax → turn scores into weights (probabilities)
3. Multiply weights by Values → get the context-aware representation

In short, each word (student) decides what other words (students) are worthy of **attention** based on how well their Q matches the other's K, and then "listens" to the most relevant ones to gather information (V).

Beyond Self-Attention: Other Key Ideas

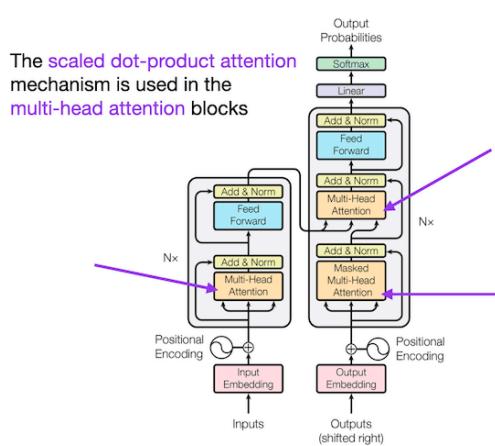
The Transformer's success comes from several innovations working together:

Innovation	What It Does
Multi-head attention	Multiple attention "views" in parallel — each head learns different patterns
Positional encoding	Adds word-order information (since attention has no inherent order)
Layer normalization	Stabilizes training by normalizing activations
Residual connections	Skip connections that prevent vanishing gradients
Feed-forward layers	Adds non-linear transformations after attention

All of these are stacked into **layers** (GPT-3 has 96 layers, GPT-4 likely has more).

Transformers

Transformers are neural network architectures introduced in the 2017 paper “Attention is All You Need”. They process sequences (like sentences) without using RNNs or convolutions—instead relying entirely on attention.



Source: "Attention Is All You Need" (<https://arxiv.org/abs/1706.03762>)

Attention's Role in Transformers:

- Self-attention lets every word in a sentence “look at” every other word to gather context.
 - Example: “Alice eats pizza” → “eats” attends to “Alice” (subject) and “pizza” (object).
- Multi-head attention runs this 8-16 times in parallel, capturing different relationships (syntax, semantics, etc.).
- Cross-attention (in decoders) lets output words attend to encoder outputs (e.g., translation)

Why Attention Defines Transformers

- Replaces recurrence: No sequential processing— everything parallelizes perfectly on GPUs.
- Captures long-range dependencies: “The key to success is... 1000 words later ...patience.”
- Dynamic focus: Attention weights adapt per input, unlike fixed convolutions.

Bottom line:

1. Transformers = stacked layers of attention + feed-forward + normalization.
2. Attention handles “who relates to whom”; feed-forward adds non-linearity.

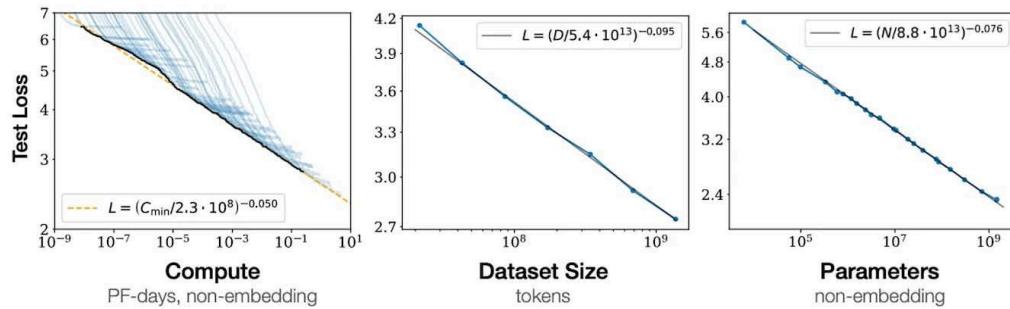
The Scaling Law for Algorithms

Another key discovery: **bigger Transformers perform better, predictably.**

As you scale up:

- Number of **parameters** (weights in the model)
- Amount of **training data**
- Amount of **compute**

The model's loss decreases following a **smooth power law**.

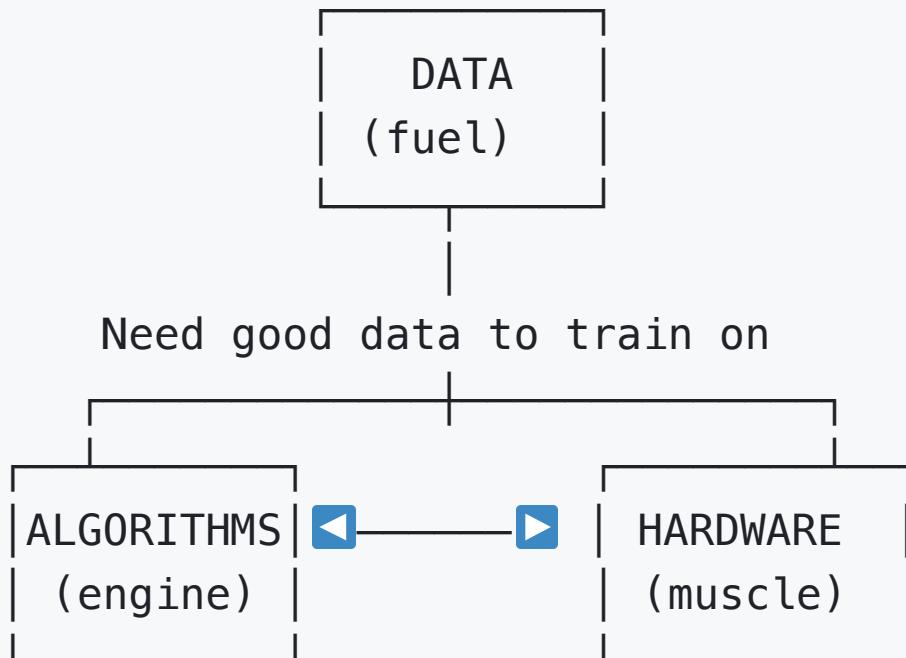


This gave researchers confidence to invest billions in training ever-larger models, because they could **predict performance in advance**.

Model	Parameters	Year
GPT-2	1.5 billion	2019
GPT-3	175 billion	2020
GPT-4	~1.8 trillion (rumored MoE)	2023

How the Three Pillars Interact

The Interplay – No Pillar Works Alone



Better algorithms
need hardware to enable bigger
run at scale algorithms & data

Each pillar **amplifies** the other two. Breakthroughs in one pillar often unlock progress in the others.

This is why:

- Companies are trying to buy hardware (NVIDIA and fast DRAM)
- Companies are investing in data partnerships and synthetic data generation
- Researchers are inventing new architectures (MoE, long context) to better utilize data and

Historical Timeline — Convergence of the Three Pillars

Year	Data	Algorithm	Hardware
2012	ImageNet (1M images)	AlexNet (CNN)	NVIDIA GTX 580
2017	Large web corpora	Transformer	V100 GPUs
2018	BookCorpus + Wikipedia	BERT	TPU v3 pods
2020	300B tokens (web)	GPT-3 (175B params)	Thousands of V100s
2022	Chinchilla-optimal data	Instruction tuning, RLHF	A100 clusters
2023	2T+ tokens, curated	GPT-4, Claude	H100 clusters
2024	15T+ tokens, synthetic	MoE, long context	H200, B200

Notice how **all three pillars advance together** — no single breakthrough alone was sufficient.

The Cost Equation

Training a state-of-the-art LLM requires investment in all three pillars:

Pillar	Cost Component	Example
Data	Collection, cleaning, licensing	Web scraping infra, data partnerships
Algorithms	Research teams, experimentation	Hundreds of researchers
Hardware	GPU clusters, electricity, cooling	10,000+ H100s at \$30k each

Estimated total cost to train a frontier model (2024):
\$500M–\$1B+

What This Means for You as a Developer

You don't need to train your own LLM. But understanding the three pillars helps you:

Data:

- Understand why RAG and fine-tuning work (you're adding better data)
- Know that data quality directly affects output quality

Algorithms:

- Understand why Transformers power modern AI
- Appreciate trade-offs between model size and capability

Hardware:

- Choose the right hardware for running local models
- Understand why some models are fast (small, quantized) vs slow (large, full precision)
- Know why cloud GPU costs matter for AI projects

Important Architectural Variants

Not all LLMs use the same Transformer design:

Variant	Architecture	Examples
Decoder-only	Predicts next token (autoregressive)	GPT, Claude, LLaMA
Encoder-only	Understands input (bidirectional)	BERT, RoBERTa
Encoder-decoder	Input → output (sequence-to-sequence)	T5, BART
Mixture of Experts (MoE)	Only activates a subset of parameters per token	Mixtral, GPT-4 (rumored)

Most modern chat LLMs use **decoder-only** architecture
— they generate text one token at a time, left to right.

Summary

Pillar	Key Takeaway
Data	Trillions of high-quality tokens; quality > quantity; data wall is approaching
Algorithms	Transformer + self-attention was the breakthrough; scaling laws guide progress
Hardware	GPUs enabled parallel training; NVIDIA dominates; quantization democratizes access

The LLM revolution happened because all three pillars matured simultaneously.

None alone was sufficient — their **convergence** created the AI moment we are living through.

Discussion Questions

1. If high-quality text data runs out, what alternatives could sustain LLM improvement?
2. Could a fundamentally new architecture replace the Transformer? What might it look like?
3. As hardware costs drop, will every company eventually train its own LLM, or will the "API economy" dominate?
4. Which of the three pillars do you think will be the **biggest bottleneck** in the next 5 years?